

STATS3850 - Assignment #1

Question 1

2.3 Q9. This exercise involves the Auto data set studied in the lab. Make sure that the missing values have been removed from the data.

```
library(ISLR)
dim(Auto)
```

```
## [1] 392 9
```

```
summary(Auto)
```

```
##      mpg      cylinders  displacement  horsepower
##  Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0
## 1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0
## Median :22.75   Median :4.000   Median :151.0   Median : 93.5
## Mean   :23.45   Mean   :5.472   Mean   :194.4   Mean   :104.5
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0
## Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0
##
##      weight      acceleration      year      origin
##  Min.   :1613   Min.   : 8.00   Min.   :70.00   Min.   :1.000
## 1st Qu.:2225   1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000
## Median :2804   Median :15.50   Median :76.00   Median :1.000
## Mean   :2978   Mean   :15.54   Mean   :75.98   Mean   :1.577
## 3rd Qu.:3615   3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000
## Max.   :5140   Max.   :24.80   Max.   :82.00   Max.   :3.000
##
##      name
## amc matador      : 5
## ford pinto       : 5
## toyota corolla    : 5
## amc gremlin       : 4
## amc hornet        : 4
## chevrolet chevette: 4
## (Other)          :365
```

- (a) Which of the predictors are quantitative, and which are qualitative? **quantitative:** mpg, cylinders, displacement, horsepower, weight, # acceleration, year **qualitative:** name, origin
- (b) What is the range of each quantitative predictor? You can answer this using the range() function.

```
sapply(Auto[, 1:7], range)
```

```
##      mpg cylinders displacement horsepower weight acceleration year
## [1,]  9.0         3           68          46    1613           8.0   70
## [2,] 46.6         8          455         230    5140          24.8   82
```

- (c) What is the mean and standard deviation of each quantitative predictor?

```
print("Means")
```

```
## [1] "Means"
```

```
sapply(Auto[, 1:7], mean)
```

```
##          mpg      cylinders displacement  horsepower      weight
## 23.445918    5.471939    194.411990    104.469388  2977.584184
## acceleration      year
## 15.541327    75.979592
```

```
print("Standard Deviations")
```

```
## [1] "Standard Deviations"
```

```
sapply(Auto[, 1:7], sd)
```

```
##          mpg      cylinders displacement  horsepower      weight
##  7.805007    1.705783    104.644004    38.491160    849.402560
## acceleration      year
##  2.758864    3.683737
```

- (d) Now remove the 10th through 85th observations. What is the range, mean, and standard deviation of each predictor in the subset of the data that remains?

```
newAuto = Auto[-(10:85),]
```

```
# stats
```

```
sapply(newAuto[, 1:7], range)
```

```
##          mpg cylinders displacement horsepower weight acceleration year
## [1,] 11.0          3           68          46    1649           8.5    70
## [2,] 46.6          8          455          230    4997          24.8    82
```

```
sapply(newAuto[, 1:7], mean)
```

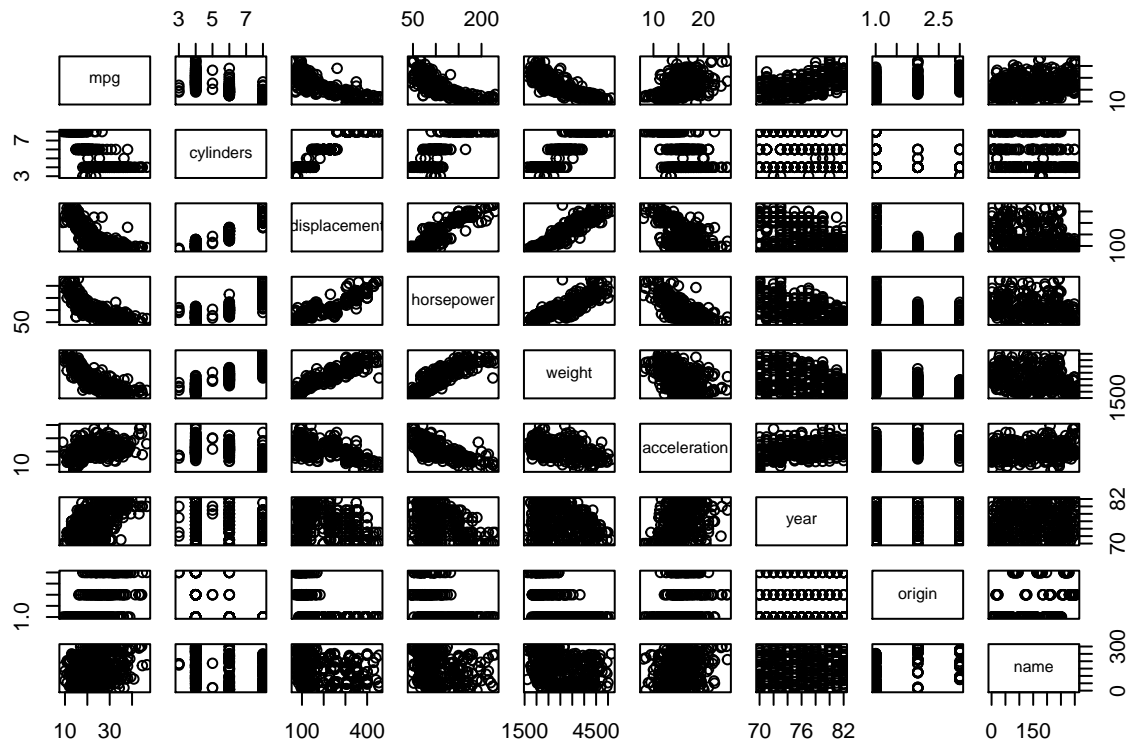
```
##          mpg      cylinders displacement  horsepower      weight
## 24.404430    5.373418    187.240506    100.721519  2935.971519
## acceleration      year
## 15.726899    77.145570
```

```
sapply(newAuto[, 1:7], sd)
```


```
##          mpg      cylinders displacement  horsepower      weight
##  7.867283    1.654179    99.678367    35.708853    811.300208
## acceleration      year
##  2.693721    3.106217
```

- (e) Using the full data set, investigate the predictors graphically, using scatterplots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings.


```
pairs(Auto)
```




```
plot(Auto$mpg, Auto$weight)
```



```
# The heavier the car the lower the mpg.
plot(Auto$mpg, Auto$displacement)
```



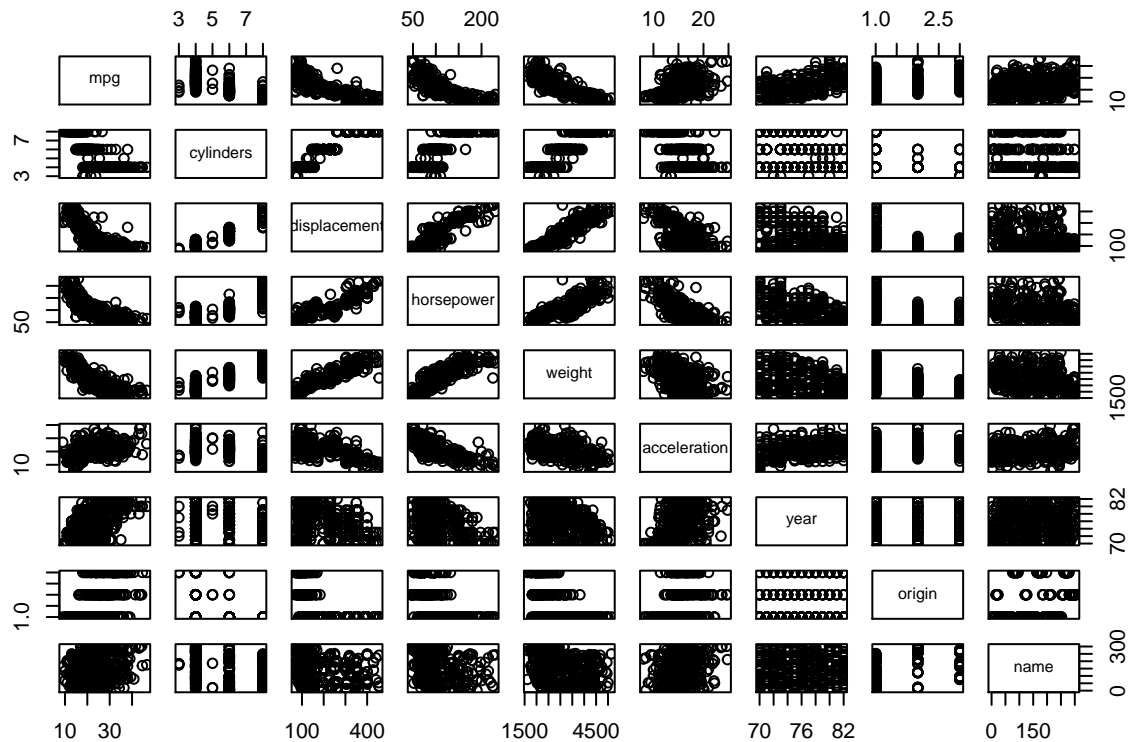
```
# Higher displacement, less Mpg.
plot(Auto$year, Auto$horsepower)
```



```
# It appears the newer cars have less horsepower than the older ones.
```

- (f) Suppose that we wish to predict gas mileage (mpg) on the basis of the other variables. Do your plots suggest that any of the other variables might be useful in predicting mpg? Justify your answer.

```
pairs(Auto)
```



It looks like weight, horsepower and displacement are clear indicators of MPG.

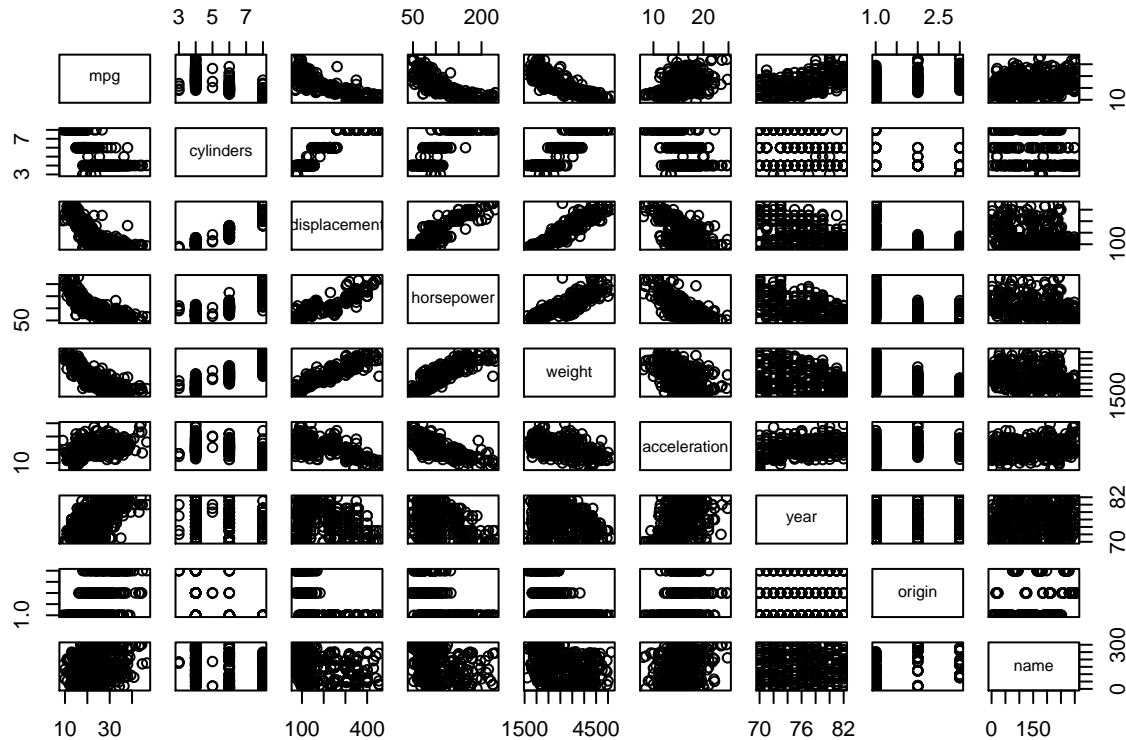
3.6 Q9. This question involves the use of multiple linear regression on the Auto data set.

```
library(ISLR)
summary(Auto)
```

```
##      mpg      cylinders  displacement  horsepower
##  Min.   : 9.00    Min.   :3.000    Min.   : 68.0    Min.   : 46.0
## 1st Qu.:17.00    1st Qu.:4.000    1st Qu.:105.0    1st Qu.: 75.0
## Median :22.75    Median :4.000    Median :151.0    Median : 93.5
## Mean   :23.45    Mean   :5.472    Mean   :194.4    Mean   :104.5
## 3rd Qu.:29.00    3rd Qu.:8.000    3rd Qu.:275.8    3rd Qu.:126.0
## Max.   :46.60    Max.   :8.000    Max.   :455.0    Max.   :230.0
##
##      weight  acceleration      year      origin
##  Min.   :1613    Min.   : 8.00    Min.   :70.00    Min.   :1.000
## 1st Qu.:2225    1st Qu.:13.78    1st Qu.:73.00    1st Qu.:1.000
## Median :2804    Median :15.50    Median :76.00    Median :1.000
## Mean   :2978    Mean   :15.54    Mean   :75.98    Mean   :1.577
## 3rd Qu.:3615    3rd Qu.:17.02    3rd Qu.:79.00    3rd Qu.:2.000
## Max.   :5140    Max.   :24.80    Max.   :82.00    Max.   :3.000
##
##      name
## amc matador      : 5
## ford pinto       : 5
## toyota corolla   : 5
## amc gremlin      : 4
## amc hornet       : 4
## chevrolet chevette: 4
## (Other)          :365
```

(a) Produce a scatterplot matrix which includes all of the variables in the data set.

```
pairs(Auto)
```



(b) Compute the matrix of correlations between the variables using the function `cor()`. You will need to exclude the name variable, which is qualitative.

```
cor(subset(Auto, select=name))
```

```
##           mpg  cylinders displacement horsepower  weight
## mpg      1.000000 -0.7776175  -0.8051269 -0.7784268 -0.8322442
## cylinders -0.7776175  1.0000000   0.9508233  0.8429834  0.8975273
## displacement -0.8051269  0.9508233   1.0000000  0.8972570  0.9329944
## horsepower  -0.7784268  0.8429834   0.8972570  1.0000000  0.8645377
## weight     -0.8322442  0.8975273   0.9329944  0.8645377  1.0000000
## acceleration  0.4233285 -0.5046834  -0.5438005 -0.6891955 -0.4168392
## year        0.5805410 -0.3456474  -0.3698552 -0.4163615 -0.3091199
## origin      0.5652088 -0.5689316  -0.6145351 -0.4551715 -0.5850054
##
## acceleration  year  origin
## mpg      0.4233285  0.5805410  0.5652088
## cylinders -0.5046834 -0.3456474 -0.5689316
## displacement -0.5438005 -0.3698552 -0.6145351
## horsepower  -0.6891955 -0.4163615 -0.4551715
## weight     -0.4168392 -0.3091199 -0.5850054
## acceleration  1.0000000  0.2903161  0.2127458
## year        0.2903161  1.0000000  0.1815277
## origin      0.2127458  0.1815277  1.0000000
```

(c) Use the `lm()` function to perform a multiple linear regression with `mpg` as the response and all other variables except `name` as the predictors. Use the `summary()` function to print the results. Comment on the output. For instance:

i. Is there a relationship between the predictors and the response? Yes, there is clearly a relationship

between these variables and the response. This is evident by the p-values being significant and the handful of co-efficients not being ~ 0 .

- ii. Which predictors appear to have a statistically significant relationship to the response? Displacement, weight, year, and origin. Judged by p-values of each predictors t-value.
- iii. What does the coefficient for the year variable suggest? It's **year's** co-efficient of 0.7508 seems to suggest that for ever year brings a increase 0.75 mpg increase.



```
fit_1 = lm(mpg~.-name, data=Auto)
summary(fit_1)



##
## Call:
## lm(formula = mpg ~ . - name, data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.5903 -2.1565 -0.1169  1.8690 13.0604
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -17.218435   4.644294  -3.707  0.00024 ***
## cylinders     -0.493376   0.323282  -1.526  0.12780
## displacement  0.019896   0.007515   2.647  0.00844 **
## horsepower    -0.016951   0.013787  -1.230  0.21963
## weight        -0.006474   0.000652  -9.929 < 2e-16 ***
## acceleration  0.080576   0.098845   0.815  0.41548
## year           0.750773   0.050973  14.729 < 2e-16 ***
## origin         1.426141   0.278136   5.127 4.67e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.328 on 384 degrees of freedom
## Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
## F-statistic: 252.4 on 7 and 384 DF,  p-value: < 2.2e-16
```

Answer: i. Is there a relationship between the predictors and the response? * Yes, there is a relationship between the predictors and the response by testing the null hypothesis of whether all the regression coefficients are zero. The F -statistic is far from 1 (with a small p-value), indicating evidence against the null hypothesis.

- ii. Which predictors appear to have a statistically significant relationship to the response?
 - Looking at the p-values associated with each predictor's t-statistic, we see that displacement, weight, year, and origin have a statistically significant relationship, while cylinders, horsepower, and acceleration do not.
- iii. What does the coefficient for the year variable suggest?
 - The regression coefficient for year, 0.7508, suggests that for every one year, mpg increases by the coefficient. In other words, cars become more fuel efficient every year by almost 1 mpg / year.
- (d) Use the plot() function to produce diagnostic plots of the linear regression fit. Comment on any problems you see with the fit. Do the residual plots suggest any unusually large outliers? Does the leverage plot identify any observations with unusually high leverage?

```
plot(fit_1)
```





```
#plot(predict(fit_1), rstudent(fit_1))
```

```
# The residual plot shows that the model may not be a very good fit. Since it follows a curve and as op
```

```
# Point 14 seems to have unusually high leverage
```

- (e) Use the * and : symbols to fit linear regression models with interaction effects. Do any interactions appear to be statistically significant? I used the co-relation matrix to find the most co-related variables. They tend to be the variables with interaction effects.

```
fit_2 = lm(mpg~cylinders*displacement+displacement*weight, data=Auto)
summary(fit_2)
```

```
##
## Call:
## lm(formula = mpg ~ cylinders * displacement + displacement *
##     weight, data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.2934  -2.5184  -0.3476   1.8399  17.7723
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    5.262e+01  2.237e+00  23.519  < 2e-16 ***
## cylinders      7.606e-01  7.669e-01   0.992   0.322
## displacement  -7.351e-02  1.669e-02  -4.403  1.38e-05 ***
## weight        -9.888e-03  1.329e-03  -7.438  6.69e-13 ***
## cylinders:displacement -2.986e-03  3.426e-03  -0.872   0.384
## displacement:weight   2.128e-05  5.002e-06   4.254  2.64e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.103 on 386 degrees of freedom
## Multiple R-squared:  0.7272, Adjusted R-squared:  0.7237
## F-statistic: 205.8 on 5 and 386 DF,  p-value: < 2.2e-16
```

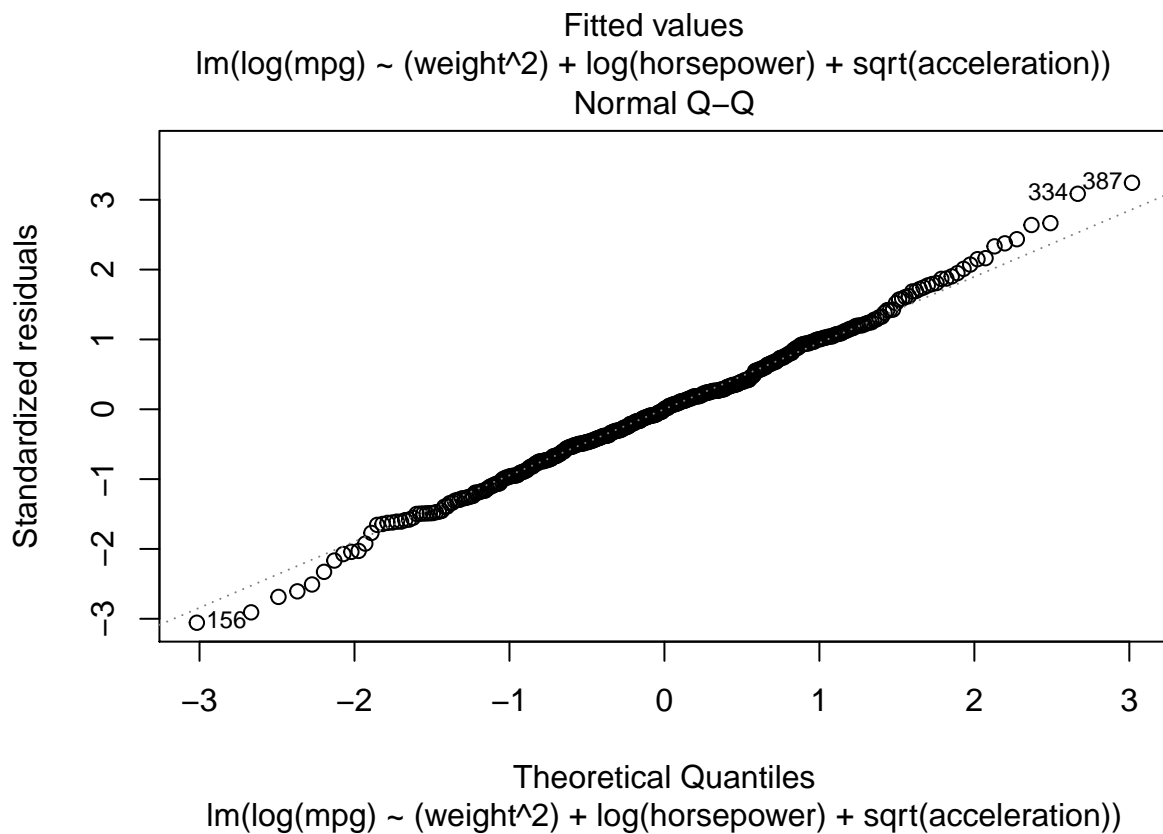

- (f) Try a few different transformations of the variables, such as log(X), root(X), X². Comment on your findings.

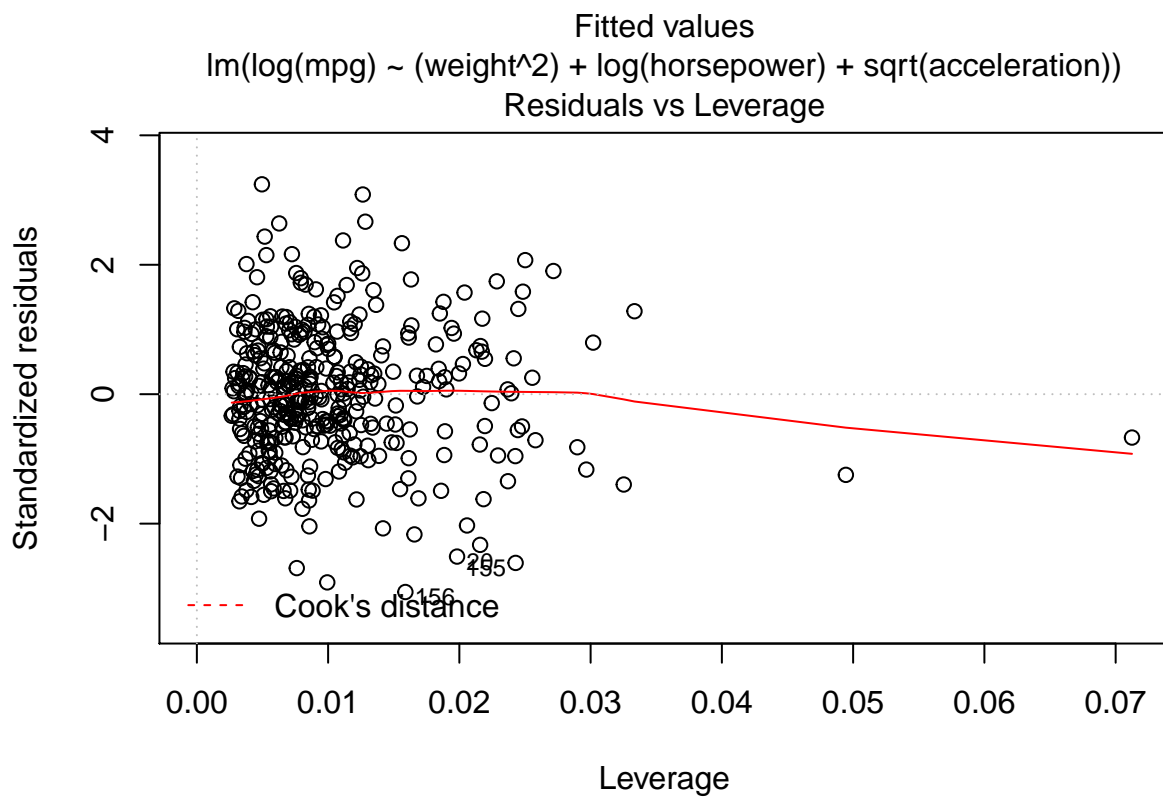

```
fit_3 = lm(log(mpg)~(weight^2)+log(horsepower)+sqrt(acceleration), data=Auto)
summary(fit_3)
```

```
##
## Call:
## lm(formula = log(mpg) ~ (weight^2) + log(horsepower) + sqrt(acceleration),
##     data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.46240 -0.09677  0.00029  0.09750  0.49334
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.422e+00  3.997e-01  16.069  < 2e-16 ***
## weight        -1.889e-04  2.282e-05  -8.281  2.00e-15 ***
## log(horsepower) -5.100e-01  7.225e-02  -7.058  7.84e-12 ***
## sqrt(acceleration) -1.073e-01  3.772e-02  -2.845  0.00467 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1525 on 388 degrees of freedom
```

```
## Multiple R-squared:  0.8004, Adjusted R-squared:  0.7988  
## F-statistic: 518.6 on 3 and 388 DF,  p-value: < 2.2e-16
```

```
plot(fit_3)
```





This set of transformations actually seems to help the model achieve a nearly straight residual curve

Question 2

In this exercise you will create some simulated data and will fit simple linear regression models to it. Make sure to use `set.seed(1)` prior to starting part (a) to ensure consistent results.

```
set.seed(1)
```

- (a) Using the `rnorm()` function, create a vector, `x`, containing 100 observations drawn from a $N(0,1)$ distribution. This represents a feature, X .

```
x <- rnorm(100)
```

- (b) Using the `rnorm()` function, create a vector, `eps`, containing 100 observations drawn from a $N(0,0.25)$ distribution i.e. a normal distribution with mean zero and variance 0.25.

```
eps <- rnorm(100, 0, 0.25)
```


- (c) Using `x` and `eps`, generate a vector `y` according to the model. What is the length of the vector `y`? What are the values of B_0 and B_1 in this linear model?

Clearly, $B_0 = -1$ and $B_1 = 0.5$

```
y <- (0.5 * x) + eps - 1
```

- (d) Create a scatterplot displaying the relationship between `x` and `y`. Comment on what you observe.

```
plot(x,y)
```



- (e) Fit a least squares linear model to predict `y` using `x`. Comment on the model obtained. How do \hat{B}_0 and \hat{B}_1 compare to B_0 and B_1 ? $\hat{B}_0 = -1.00942$ & $\hat{B}_1 = 0.49973$, The estimates are quite close to the actual coefficients of the population model.


```
fit_3 <- lm(y~x)
summary(fit_3)
```

```
##
## Call:
## lm(formula = y ~ x)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.46921 -0.15344 -0.03487  0.13485  0.58654
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.00942    0.02425  -41.63  <2e-16 ***
## x             0.49973    0.02693   18.56  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2407 on 98 degrees of freedom
## Multiple R-squared:  0.7784, Adjusted R-squared:  0.7762
## F-statistic: 344.3 on 1 and 98 DF,  p-value: < 2.2e-16
```

- (f) Display the least squares line on the scatterplot obtained in (d). Draw the population regression line on the plot, in a different color. Use the legend() command to create an appropriate legend.

```
plot(x,y,pch=16, xlab="X",ylab="Y",main="X vs. Y")
abline(coefficients(fit_3), lwd=2, col=2,lty=1)
abline(a=-1,b=0.5,col=3,lty=2)
legend( x= "topleft",
       legend=c("Model","Regression line"),
       col=c(2, 3), lty=1)
```



- (g) Now fit a polynomial regression model that predicts y using x and x^2 . Is there evidence that the quadratic term improves the model fit? Explain your answer. EDIT: There is evidence that model fit has increased over the training data given the slight increase in R^2 and RSE. Although, the p-value of the t-statistic suggests that there isn't a relationship between y and x^2 .

```
model_squared_x = lm(y~x+I(x^2))
summary(model_squared_x)
```

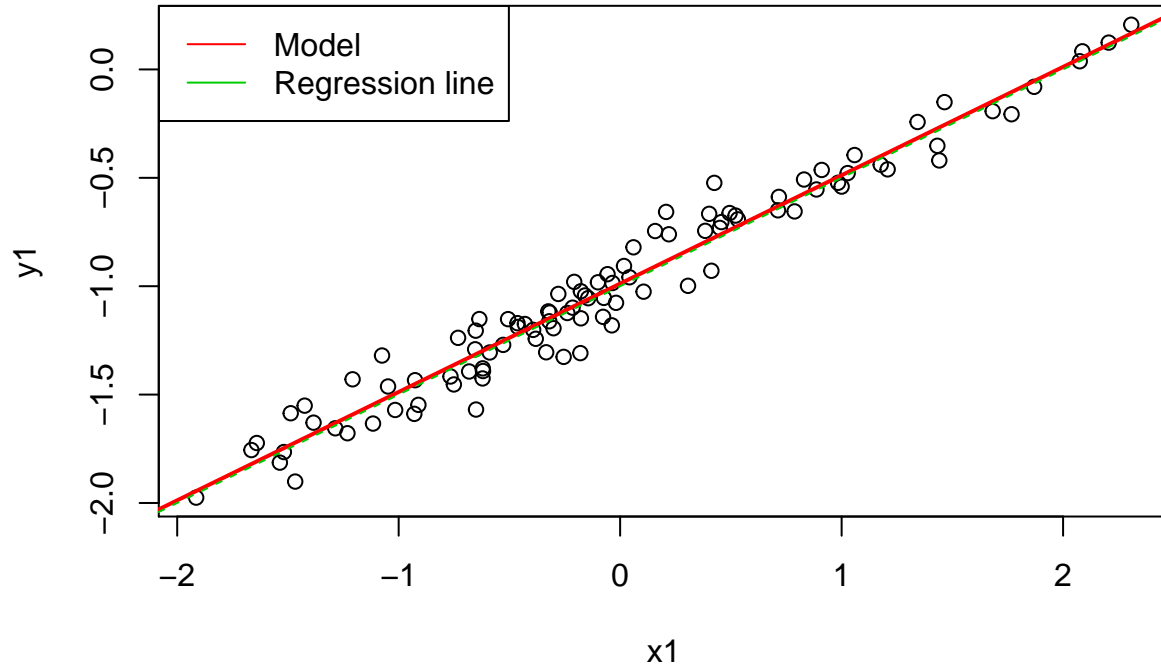
```
##
## Call:
## lm(formula = y ~ x + I(x^2))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.4913 -0.1563 -0.0322  0.1451  0.5675
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.98582    0.02941  -33.516  <2e-16 ***
## x            0.50429    0.02700   18.680  <2e-16 ***
## I(x^2)       -0.02973    0.02119   -1.403    0.164
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2395 on 97 degrees of freedom
## Multiple R-squared:  0.7828, Adjusted R-squared:  0.7784
## F-statistic: 174.8 on 2 and 97 DF,  p-value: < 2.2e-16
```

- (h) Repeat (a)–(f) after modifying the data generation process in such a way that there is less noise in the data. The model (3.39) should remain the same. You can do this by decreasing the variance of the normal distribution used to generate the error term in (b). Describe your results.

```
set.seed(1)
eps_2 = rnorm(100, 0, 0.11)
x1 = rnorm(100)
y1 = -1 + 0.5*x1 + eps_2
plot(x1, y1)
fit_2h = lm(y1~x1)
summary(fit_2h)
```

```
##
## Call:
## lm(formula = y1 ~ x1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.255657 -0.066397  0.000589  0.064135  0.252248
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.988026    0.009938  -99.42  <2e-16 ***
## x1           0.499897    0.010419   47.98  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0993 on 98 degrees of freedom
## Multiple R-squared:  0.9592, Adjusted R-squared:  0.9587
## F-statistic: 2302 on 1 and 98 DF,  p-value: < 2.2e-16
```

```
abline(coefficients(fit_2h), lwd=2, col=2,lty=1)
abline(a=-1,b=0.5,col=3,lty=2)
legend( x= "topleft",
       legend=c("Model","Regression line"),
       col=c(2, 3), lty=1)
```



- (i) Repeat (a)–(f) after modifying the data generation process in such a way that there is more noise in the data. The model (3.39) should remain the same. You can do this by increasing the variance of the normal distribution used to generate the error term in (b). Describe your results.

```
set.seed(1)
eps_3 = rnorm(100, 0, 1)
x1 = rnorm(100)
y1 = -1 + 0.5*x1 + eps_3
plot(x1, y1)
fit_2i = lm(y1~x1)
summary(fit_2i)
```

```
##
## Call:
## lm(formula = y1 ~ x1)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.32416	-0.60361	0.00536	0.58305	2.29316

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.89115	0.09035	-9.864	2.39e-16 ***
x1	0.49907	0.09472	5.269	8.16e-07 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```



```
## Residual standard error: 0.9028 on 98 degrees of freedom
## Multiple R-squared:  0.2207, Adjusted R-squared:  0.2128
## F-statistic: 27.76 on 1 and 98 DF,  p-value: 8.158e-07
```

```
abline(coefficients(fit_2i), lwd=2, col=2,lty=1)
abline(a=-1,b=0.5,col=3,lty=2)
legend( x= "topleft",
       legend=c("Model","Regression line"),
       col=c(2, 3), lty=1)
```



- (j) What are the confidence intervals for B_0 and B_1 based on the original data set, the noisier data set, and the less noisy data set? Comment on your results.

```
confint(fit_3)
```

```
##                2.5 %    97.5 %
## (Intercept) -1.0575402 -0.9613061
## x           0.4462897  0.5531801
```

```
confint(fit_2h)
```

```
##                2.5 %    97.5 %
## (Intercept) -1.0077485 -0.9683041
## x1          0.4792205  0.5205743
```

```
confint(fit_2i)
```

```
##                2.5 %    97.5 %
## (Intercept) -1.0704405 -0.7118552
## x1          0.3110958  0.6870395
```

Question 3

a)

```

n=10;
x<-rnorm(n);
y<-4.20+6.9*x+rnorm(n,0,2.5)
fit_1<-lm(y~x)
print("Beta 0 & 1")

## [1] "Beta 0 & 1"

summary(fit_1)$coefficients[1:2,1]

## (Intercept)          x
##    4.115329      7.074350

```

```

b)

vcov(fit_1)

##              (Intercept)          x
## (Intercept)  0.6952192 -0.2024136
## x           -0.2024136  0.3839785

```

```

c)

beta_0_list <- c(1:1000)
beta_1_list <- c(1:1000)
for (i in 1:1000){
  n=10
  x<-rnorm(n)
  y<-4.20+6.9*x+rnorm(n,0,2.5)
  fit_2<-lm(y~x)
  beta_0_list[i] <- summary(fit_2)$coefficients[1,1]
  beta_1_list[i] <- summary(fit_2)$coefficients[2,1]
}
print('Covariance')

```

```

## [1] "Covariance"

cov(beta_0_list, beta_1_list)

## [1] -0.0003285326

```

Question 4

```

# 1

q_4_func <- function(n) {
  x<-rnorm(n);
  y<-4.20+6.9*x+rnorm(n,0,2.5)
  fit_2<-lm(y~x)
  t0 <- summary(fit_2)$coefficients[2,3]
  return(t0)
}

t0 <- q_4_func(100)
cat("The value of t0: ", t0)

## The value of t0: 24.85661

```

```

# 2
values = list()
for (i in 1:1000) {
  values[i] <- q_4_func(100)
}
count = 0
for (k in values){
  if(k > t0){
    count = count + 1
  }
}
portion <- count/1000
cat("Portion of ti > t0: ", portion)

```

```
## Portion of ti > t0: 0.837
```